
DD2424 - Assignment 1 (Bonus)

SUMMARY

I have completed both parts of the bonus assignment exercises. For the first part, I have used the entirety of the training data, implemented training data augmentation, and used a decaying learning rate. I also experimented with a number of different values for the regularization parameter. For the second part, I have derived the gradients for the multiple binary cross-entropy loss and compared the results when using the sigmoid activation with the corresponding results for the softmax.

Oskar STIGLAND
DD2424
Spring 2023

Part I - Improving Performance of the Network

For this part, I have experimented with different parameter setting while adding the improvements. I recorded the best results and the (seemingly) most stable training procedure with the following parameter settings:

- $\text{batchN} = 50$
- $\text{epochsN} = 50$
- $\lambda = 0.1$
- $\eta = 0.01$
- $\eta_n = 10$
- $p_{\text{flip}} = 0.2$

where the η_n corresponds to the number of epochs per decay of the learning rate. Per the assignment suggestion, I initially used a horizontal flip probability of $p = 0.5$ for the training data, performing a shuffle and flip for every epoch. However, it seems a lower flip probability works equally well and I thus settled for $p = 0.2$. As mentioned, I have also utilized all of the training data, reducing the validation data to a smaller subset of the training data with $N_{\text{val}} = 1000$. I also experimented with a number of different settings for the regularization parameter. It would seem $\lambda = 0.1$ is a fairly optimal value for training. For the above parameters, the model reaches an accuracy of 40.76% and across the different improvements it seems that expanding and augmenting the dataset yields the largest benefits in terms of accuracy, i.e. by using the entirety of the dataset with a non-zero flip probability.




Figure 1: Loss and accuracy for optimal set of parameters




Figure 2: Class-specific weights for optimal set of parameters

Part II - Multiple Binary Cross-Entropy Losses

Obtaining the gradient

In order to obtain the expression for $\partial\ell/\partial\mathbf{s}$, we consider the chain rule, where

$$\frac{\partial\ell}{\partial\mathbf{s}} = \frac{\partial\ell}{\partial\mathbf{p}} \frac{\partial\mathbf{p}}{\partial\mathbf{s}}$$

First, we consider differentiating the class-specific loss w.r.t. the corresponding probability, s.t. $\ell = -1/K \sum_{k=1}^K \ell_k$:

$$\frac{\partial\ell_k}{\partial p_k} = \frac{\partial}{\partial p_k} [(1 - y_k) \log(1 - p_k) + y_k \log(p_k)] = \frac{y_k - 1}{1 - p_k} + \frac{y_k}{p_k}$$

Second, we recognize that $p_k = \sigma(s_k)$, where $\sigma : \mathbb{R} \rightarrow [0, 1]$ is the logistic sigmoid function, for which we also have that

$$\sigma(s) = \frac{1}{1 + e^{-s}} = \frac{e^s}{e^s + 1}$$

as well as $1 - \sigma(s) = \sigma(-s)$ and $\sigma'(s) = \sigma(s)(1 - \sigma(s))$. Hence, we get that

$$\begin{aligned} \frac{\partial\ell_k}{\partial s_k} &= \frac{\partial}{\partial s_k} [(1 - y_k) \log(1 - \sigma(s_k)) + y_k \log(\sigma(s_k))] \\ &= \frac{y_k - 1}{1 - \sigma(s_k)} \sigma(s_k)(1 - \sigma(s_k)) + \frac{y_k}{\sigma(s_k)} \sigma(s_k)(1 - \sigma(s_k)) \\ &= (y_k - 1) \sigma(s_k) + y_k (1 - \sigma(s_k)) \\ &= y_k - \sigma(s_k) \end{aligned}$$

Hence, the Jacobians we're looking for are given by

$$\begin{aligned} \frac{\partial\ell}{\partial\mathbf{p}} &= -\frac{1}{K} [\mathbf{y}^T \text{diag}(\mathbf{p})^{-1} - (\mathbf{1} - \mathbf{y})^T \text{diag}(\mathbf{1} - \mathbf{p})^{-1}] \\ \frac{\partial\mathbf{p}}{\partial\mathbf{s}} &= \text{diag}(\mathbf{p}) \text{diag}(\mathbf{1} - \mathbf{p}) \end{aligned}$$

such that

$$\begin{aligned} \frac{\partial\ell}{\partial\mathbf{s}} &= -\frac{1}{K} \left[\mathbf{y}^T \underbrace{\text{diag}(\mathbf{p})^{-1} \text{diag}(\mathbf{p}) \text{diag}(\mathbf{1} - \mathbf{p})}_{=\text{diag}(\mathbf{1} - \mathbf{p})} - (\mathbf{1} - \mathbf{y})^T \underbrace{\text{diag}(\mathbf{1} - \mathbf{p})^{-1} \text{diag}(\mathbf{p}) \text{diag}(\mathbf{1} - \mathbf{p})}_{=\text{diag}(\mathbf{p})} \right] \\ &= -\frac{1}{K} [\mathbf{y}^T \text{diag}(\mathbf{1} - \mathbf{p}) - (\mathbf{1} - \mathbf{y})^T \text{diag}(\mathbf{p})] \\ &= -\frac{1}{K} [\mathbf{y} - \mathbf{y}^T \text{diag}(\mathbf{p}) - \mathbf{p} + \mathbf{y}^T \text{diag}(\mathbf{p})] \\ &= -\frac{1}{K} (\mathbf{y} - \mathbf{p}) \end{aligned}$$

Hence, using the result from the previous loss, we should have that

$$\frac{\partial\ell}{\partial\mathbf{W}} = -\frac{1}{K} (\mathbf{y} - \mathbf{p}) \mathbf{x}^T, \quad \frac{\partial\ell}{\partial\mathbf{b}} = -\frac{1}{K} (\mathbf{y} - \mathbf{p})$$

Results

I implemented the linear classifier with a **sigmoid** activation and the multiple binary cross-entropy loss for a number of parameter settings. The loss, accuracy and class-specific weights have been plotted on the following page. Due to the $1/K$ factor in the loss, I have used a slightly higher learning rate for this part, $\eta = 0.01$ or $\eta = 0.05$, with a slightly modified decay rate, e.g. reducing the learning rate by half every 10 steps. Overall, the results are comparable to those when using the **softmax** activation, but the achieved accuracy is actually lower across all tests compared to the highest accuracy achieved with the **softmax**. However, it does seem that similarly to using the **softmax**, the highest accuracy is achieved when using a moderate regularization parameter - again, in this case $\lambda = 0.1$. In fact, the performance is somewhat worse when using $\lambda = 1.0$ as compared to using $\lambda = 0.0$.

In order to compare the **softmax** and **sigmoid** activations, I trained two separate models with the optimal set of parameters found for each type. The class-specific prediction results and the parameter setting are shown in the plots and table below. The results are in fact very similar. Both models achieve a relatively low accuracy for the category 2 and 3, while they both achieve a relatively high accuracy for categories 0, 1, and 6. The **softmax** model also seems to achieve a slightly better result for category 7. The differences, however, are marginal. I wouldn't call this overfitting, since the class-specific accuracies are fairly well-distributed across the classes, with the exception of 2 to 4 - for which both models fail.

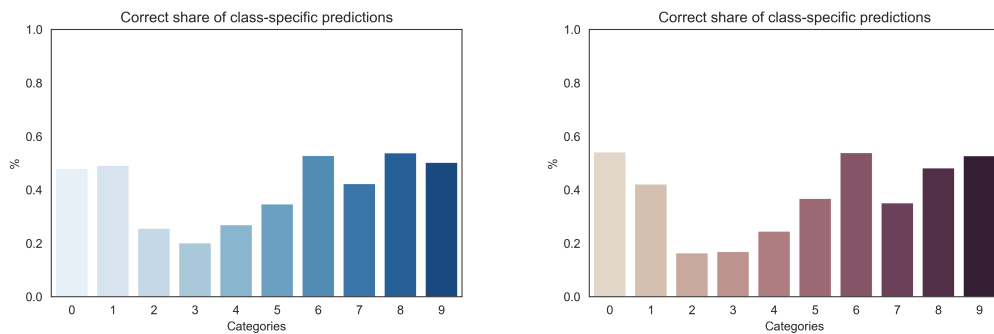
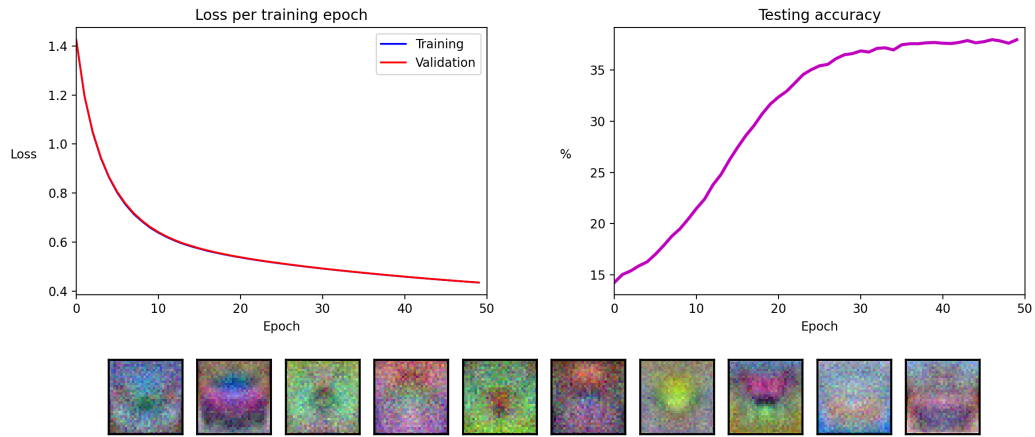
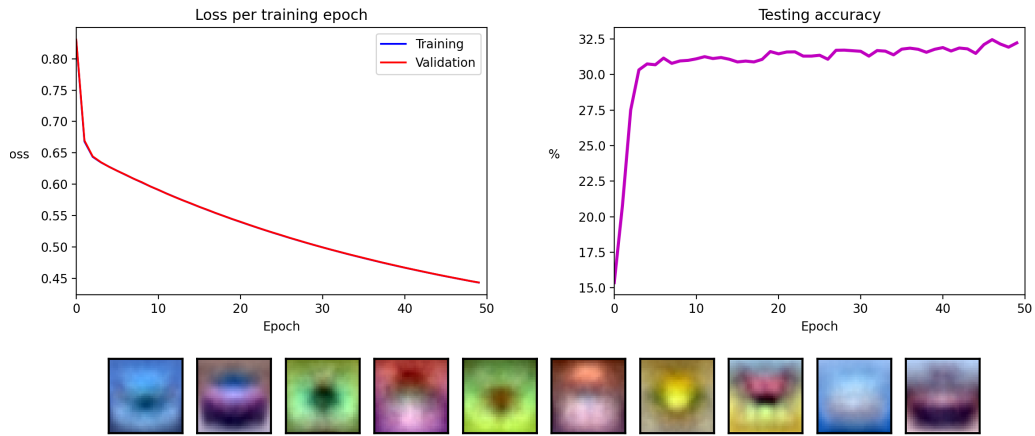
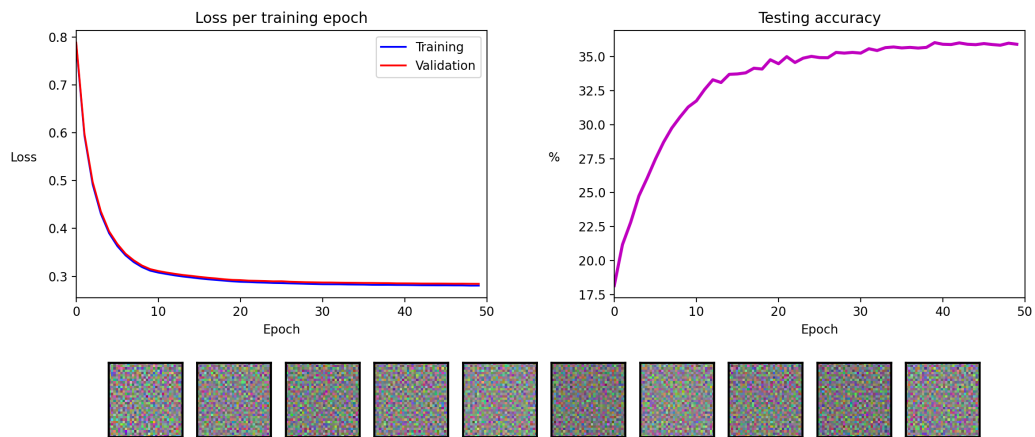


Figure 3: Class-specific accuracy for **softmax** (left) and **sigmoid** (right)

Activation	batchN	epochsN	η	η_n	λ	P_{flip}
softmax	100	50	0.01	10	0.1	0.2
sigmoid	50	50	0.01	0	0.1	0.5

Figure 4: sigmoid with $\lambda = 0.1$, $\eta = 0.01$, $\text{batchN} = 100$, $\text{epochsN} = 50$, $P_{\text{flip}} = 0.5$ Figure 5: sigmoid with $\lambda = 1.0$, $\eta = 0.01$, $\text{batchN} = 100$, $\text{epochsN} = 50$, $P_{\text{flip}} = 0.5$ Figure 6: sigmoid with $\lambda = 0.0$, $\eta = 0.01$, $\text{batchN} = 100$, $\text{epochsN} = 50$, $P_{\text{flip}} = 0.1$